

The Indian Buffet Process

Eli Bingham¹ Matt Dickenson²

¹University of North Carolina

²Duke University

February 10, 2014

- ➊ Introduction
- ➋ Dirichlet and Chinese Restaurant Processes
- ➌ Beta and Indian Buffet Processes
- ➍ Gibbs sampling
- ➎ Demonstration/Visualization
- ➏ Applications: Choice Behavior and Collaborative Filtering
- ➐ Extensions: Topic Models and Cascading IBP
- ➑ Discussion

Latent feature models

- Feature model: N items described by K features
- Dense feature model: every feature is present in every item, e.g. PCA
- Sparse feature model: only some features present in each item, and we can assume feature values and presence are independent:

$$\mathbf{F} = \mathbf{A} \otimes \mathbf{Z}$$
$$P(\mathbf{F}) = P(\mathbf{A})P(\mathbf{Z})$$




Figure : Griffiths and Ghahramani (2011) Figure 3

Example




Figure : Griffiths and Ghahramani (2011) Figure 7




Figure : Griffiths and Ghahramani (2011) Figure 9

Motivation

- Problem with finite latent feature model: K is fixed
- Goal: construct nonparametric prior on \mathbf{Z} so that K grows with the complexity of the dataset
- As with DPMMs, we can try to build one by taking $K \rightarrow \infty$ in a finite feature model




Figure : Griffiths and Ghahramani (2011) Figure 4

Background: Dirichlet Process

Finite version (Dirichlet distribution):

- Assignment of an object to a class is independent of all other assignments: $P(c|\theta) = \prod_{i=1}^N P(c_i|\theta) = \prod_{i=1}^N \theta_{c_i}$
- $\theta|\alpha \sim \text{Dirichlet}(\frac{\alpha}{K}, \dots, \frac{\alpha}{K})$ (if symmetric)
- $c_i|\theta \sim \text{Discrete}(\theta)$, where Discrete : Bernoulli :: Multinomial : Binomial

Integrating out θ : $P(c) = \frac{\prod_{k=1}^K \Gamma(m_k + \frac{\alpha}{K})}{\Gamma(\frac{\alpha}{K})^K} \frac{\Gamma(\alpha)}{\Gamma(N+\alpha)}$

As $K \rightarrow \infty$, we get the CRP




Figure : Griffiths and Ghahramani (2011) Figure 1

Background: Chinese Restaurant Process

- 1 N customers enter (in sequence) a restaurant with an infinite number of tables, each with infinite seating
- 2 First customer sits at first table with probability $\frac{\alpha}{\alpha} = 1$
- 3 i^{th} customer sits at the k^{th} table with probability $\frac{m_k}{i+\alpha-1}$, where m_k is the number of previous customers who sat at table k , or a new table with probability $\frac{\alpha}{i+\alpha-1}$




Figure : Griffiths and Ghahramani (2011) Figure 2

Limitation: each object (customer) can only belong to one class (table).

Finite latent feature models

The basic finite distribution on $z_{i,k}$ s:

$$\begin{aligned}\pi_k | \alpha &\sim \text{Beta}\left(\frac{\alpha}{K}, 1\right) \\ z_{i,k} | \pi_k &\sim \text{Bernoulli}(\pi_k)\end{aligned}$$

As with DPMMs, we can marginalize out latent feature presence probabilities π_k to obtain a distribution on matrices $\mathbf{Z} \in \{0, 1\}^{N \times K}$:

$$\begin{aligned}P(\mathbf{Z}) &= \prod_{k=1}^K \int \left(\prod_{i=1}^N P(z_{ik} | \pi_k) \right) P(\pi_k) d\pi_k \\ &= \prod_{k=1}^K \frac{\frac{\alpha}{K} \Gamma(m_k + \frac{\alpha}{K}) \Gamma(N - m_k + 1)}{\Gamma(N + 1 + \frac{\alpha}{K})}\end{aligned}$$

The $K \rightarrow \infty$ limit

- $\text{lof}(\mathbf{Z})$ is the matrix obtained by ordering the columns of \mathbf{Z} as N -digit binary numbers
- To define a probability over infinitely wide binary matrices using de Finetti's Theorem, we need exchangeable symmetry, so we define lof equivalence classes by modding out column order:




Figure : Griffiths and Ghahramani (2011) Figure 5

Indian Buffet Process


Indian Buffet Process:

- 1 N customers enter (in sequence) a buffet restaurant with an infinite number of dishes
- 2 First customer fills her plate with $\text{Poisson}(\alpha)$ number of dishes
- 3 i^{th} customer samples dishes in proportion to their popularity, with probability $\frac{m_k}{i}$, where m_k is the number of previous customers who sampled dish k
- 4 i^{th} customer then samples $K_1^{(i)} \sim \text{Poisson}(\frac{\alpha}{i})$ number of new dishes

Resulting probability distribution on matrices:

$$P(\mathbf{Z}) = \frac{\alpha^{K_+}}{\prod_{i=1}^N K_1^{(i)}!} \exp(\alpha H_N) \prod_{k=1}^{K_+} \frac{(N - m_k)!(m_k - 1)!}{N!}$$

mcdickenson.shinyapps.io/ibp-demo



Alternative derivation: Beta Process

The CRP is obtained by marginalizing over a Dirichlet process:

$$\begin{aligned}\theta|\alpha, G_0 &\sim \text{DP}(\alpha, G_0) \\ C|\theta &\sim \text{Multinomial}(\theta) \\ \text{CRP}(\alpha) &\sim \int P(C|\theta)P(\theta|\alpha, G_0)d\theta\end{aligned}$$

This representation follows from de Finetti's Theorem, which gives the existence of conditionally independent representations of infinite exchangeable joint distributions. The IBP is obtained by marginalizing over a Beta process:

$$\begin{aligned}\theta|\alpha, \beta, G_0 &\sim \text{BP}(\alpha, \beta, G_0) \\ z_i|\theta &\sim \text{BeP}(\theta) \\ \text{IBP}(\alpha, \beta) &\sim \int \prod_i P(z_i|\theta)P(\theta|\alpha, \beta, G_0)d\theta\end{aligned}$$

Alternative derivation: Stick-Breaking

- 1 Recursively break (an initially unit-length) stick, breaking off a $\text{Beta}(\alpha, 1)$ portion at each step
- 2 Let each portion of the “stick”, π_k represent the probability of each feature (sorted from largest to smallest)

This helps to show the relation between the Dirichlet process and the IBP. The stick-breaking construction is also useful for defining inference algorithms.

Properties of the Resulting Distribution

- The “effective” dimension $K_+ \sim \text{Poisson}(\alpha H_N)$
- The number of dishes on each customer’s plate is distributed $\text{Poisson}(\alpha)$ (by exchangeability)
- \mathbf{Z} remains sparse as $K \rightarrow \infty$: effective dimensions of \mathbf{Z} are $N \times K_+$, and the expected number of entries is $N\alpha$

Inference by Gibbs Sampling

Given some data \mathbf{X} , we want to sample from a marginal posterior

$$P(z_{i,k} = 1 | \mathbf{Z}_{-(ik)}, \mathbf{X}) \propto P(\mathbf{X} | \mathbf{Z}) P(z_{i,k} = 1 | \mathbf{Z}_{-(ik)})$$

Iterate continuously over the rows \mathbf{z}_i ($i = 1 \dots N$) of \mathbf{Z} :

- 1 For each column k of \mathbf{Z} , if $m_{-i,k} = 0$ (i.e. the rest of the column is empty) delete the column; otherwise set $z_{i,k} = 1$ with probability

$$P(\mathbf{X} | \mathbf{Z}) P(z_{i,k} = 1 | \mathbf{z}_{-i,k}) = P(\mathbf{X} | \mathbf{Z}) \frac{m_{-i,k}}{N}$$

- 2 At the end of the row, add $K_1^{(i)} \sim P(\mathbf{X} | \mathbf{Z}) \text{Poisson}(\frac{\alpha}{N})$ new columns with ones in row i

Variational Inference

Application 1: Choice Behavior

“A Choice Model with Infinitely Many Latent Features” (Görür, Jäkel, and Rasmussen, ICML 2006)

- Customers compare items (e.g. cell phones) based on the (binary) features of each; more features are better
- Number of features is potentially infinite and ordering is not important, so IBP is used
- Celebrity example: “With whom would you prefer to spend an hour of conversation?”




Figure : Görür, Jäkel, and Rasmussen (2006) Figure 3

Application 2: Collaborative Filtering

Extension 1: Topic Modeling

“The IBP Compound Dirichlet Process
and its Application to Focused Topic Modeling”
Williamson, Wang, Heller, and Blei (2010)

Stick-breaking construction:

$$\begin{aligned}\mu_k &\sim \text{Beta}(\alpha, 1) \\ \pi_k &= \prod_{j=1}^k \mu_j \\ b_{m,k} &\sim \text{Bernoulli}(\pi_k)\end{aligned}$$

Extension 1: Topic Modeling

Focused topic model:

- ① for $k = 1, 2, \dots$
 - Sample stick length π_k
 - Sample relative mass $\phi_k \sim \text{Gamma}(\gamma, 1)$
 - Draw topic distribution over words: $\beta_k \sim \text{Dirichlet}(\eta)$
- ② for $m = 1, \dots, M$
 - Sample binary vector b_m
 - Draw total number of words $n^{(m)} \sim \text{NB}(\sum_k b_{m,k} \phi_k, 1/2)$
 - Sample distribution over topics $\theta_m \sim \text{Dirichlet}(b_m \cdot \phi)$
 - For each word $w_{m,i}, i = 1, \dots, n^{(m)}$
 - ① Draw topic index $z_{m,i} \sim \text{Discrete}(\theta_m)$
 - ② Draw word $w_{m,i} \sim \text{Discrete}(\beta_{z_{m,i}})$

Extension 1: Topic Modeling

An advantage of the focused topic model is that it separates the global topic proportions from the distribution over topics within a topic. A rare topic within the corpus can be dominant within a document (e.g. baseball), and a frequent topic can be a small proportion of many documents.

Extension 2: Collaborative Filtering

Limitations of IBP:

- 1 Coupling of average number of features α and total number of features N_α (can be overcome with a two-parameter generalization)
- 2 Computationally complex, can be time-consuming




Figure : Griffiths and Ghahramani (2011) Figure 10